

Debugger Laboratory

CS 105, Spring 2018

Due on Tuesday, February 6, 2018, at 11:59 pm

Introduction and Goals

The goals of this assignment are to do some basic investigation of the X86-64 architecture and assembly language, and to begin learning how to use gdb. The lab pages have links to a quick gdb summary and to a printable gdb reference card; you can also find other information on Google.

It may be useful to know that you can get the compiler to generate the assembler source for a program by using “`gcc -S foo.c`”. You should also know that to use the debugger effectively, you will need to compile with the “`-g`” switch. In fact, you should just get in the habit of always compiling with “`-g`”; the situations where it’s undesirable are extremely unusual. (But note that `-S` and `-g` are best kept separate.) Also, it’s usually wise to compile *without* “`-O`” because that will make debugging more difficult, and debugging is nearly always more important than optimization. (That said, you can compile with `-S` and any optimization level, and in fact doing so can produce useful insights.)

Collect your answers to all of the following questions in a plain-text file named “`lab02.txt`”. Identify each section by problem number, and each answer by question number. Submit *only* your file `lab02.txt` using the form on the course web page.

Do not change either of the programs in this lab!

Problem 1—Debugging Optimized Code (17 Points)

Download the file `problem1.c` from the lab Web site. This file contains a function that has a small `while` loop, and a simple `main` that calls it. Briefly study the `loop_while` function to understand how it works (you don’t need to fully decode it; just get a clue about what’s going on).

It will also be useful to know what the `atoi` function does. Type “`man atoi`” in a terminal window to find out. (Note that the function is pronounced “ay-to-eye”.)

Finally, it will be useful to have a slight clue about `printf`. Since `printf` is quite complicated, for now we’ll just say that it prints answers, and “`%d`” means “print in decimal”. We encourage you to read more about `printf` in Kernighan & Ritchie or online (the advantage of reading in K&R is that the description there is less complex).

Compile the program with the `-g` switch and with **no** optimization: “`gcc -g -o problem1 problem1.c`”. Run `gdb problem1` and set a breakpoint in `main` (“`b main`”). Tell gdb not to debug the `atoi` function by

typing `skip atoi`. Run the program by typing `"r"` or `"run"`. The program will stop in `main`. (Ignore any warnings; they're meaningful but we'll work around them.)

(**Note:** to help you keep track of what you're supposed to doing, we have used italics to list the breakpoints you should have already set at the beginning of each step—except when they don't matter.)

1. *Existing breakpoint at main.*

Type `"c"` (or `"continue"`) to continue past the breakpoint. What happens?

2. *Existing breakpoint at main.*

Type `"bt"` (or `"backtrace"`) to get a trace of the call stack and find out how you got where you are. Take note of the numbers in the left column. Type `"up n"`, where *n* is one of those numbers, to get to `main`'s *stack frame* so that you can look at `main`'s variables. What file and line number are you on?

3. *Existing breakpoint at main.*

Usually when bad things happen in the library it's your fault, not the library's. In this case, the problem is that `main` passed a bad argument to `atoi`. There are two ways to find out what the bad argument is: look at `atoi`'s stack frame, or print the argument. Figure out how to look at `atoi`'s stack frame. Can you see the value that was passed?

4. *Existing breakpoint at main.*

The lack of information is sometimes caused by compiler optimizations, other times by minor debugger issues. In either case it's a nuisance. Rerun the program by typing `"r"` and let it stop at the breakpoint. Note that in step 2, `atoi` was called with the argument `"argv[1]"`. You can find out the value that was passed to `atoi` with the command `"print argv[1]"`. What is printed?

5. *Existing breakpoint at main.*

You should have learned in `cs62` that this number is the value of a `NULL` pointer. Like many library functions, `atoi` doesn't like `NULL` pointers. Rerun the program with an argument of 5 by typing `"r 5"`. Continue from the the breakpoint. What does the program print?

6. *Existing breakpoint at main.*

Without restarting `gdb`, type `"r"` (without any further parameters) to run the program yet again. (If you restarted `gdb`, you must first repeat Step 5.) When you get to the breakpoint, examine the variables `argc` and `argv` by using the `print` command. For example, type `"print argv[0]"`. Also try `"print argv[0]@argc"`, which is `gdb`'s notation for saying "print elements of the `argv` array starting at element 0 and continuing for `argc` elements." What is the value of `argc`? What are the elements of the `argv` array? Where did they come from, given that you didn't add anything to the `run` command?

7. *Existing breakpoint at main.*

The `step` or `s` command is a useful way to follow a program's execution one line at a time. Type `"s"`. Where do you wind up?

8. *Existing breakpoint at main.*

`gdb` always shows you the line that is about to be executed. Sometimes it's useful to see some context. Type `"list"` What lines do you see? Hit the return key. What do you see now?

9. *Existing breakpoint at main.*
Enter “s” to step to the next line. Then hit the return key three times. What do you think the return key does?
10. *Existing breakpoint at main.*
What are the values of `result`, `a`, and `b`?
11. Type “quit” to exit gdb. (You’ll have to tell it to kill the “inferior process”, which is the program you are debugging. Insulting!) Recompile the program, this time optimizing it with `-O1` (and including `-g` for debugging). Debug it, set a breakpoint at `loop_while` (not at `main`!), and run it with an argument of 10. Step three times. What four lines of code from `program1.c` are shown to you? Why do you think the debugger is showing you those lines in that order?
12. Quit gdb again and recompile with `-O2`. Debug the program. Disassemble the `main` function by typing “disassem main”. What is the address of the instruction that calls `atoi`? What is the address of the instruction that calls `printf`? (You will have to do some deduction here, because `gcc` mangles the names a bit.)
13. What is the address of the instruction that calls `loop_while`?
14. A handy feature of `print` is that you can use it to convert between bases. For example, what happens when you type “`print/x 42`”? How about “`p 0x2f`”?
15. We may not have covered it yet, but functions return results in `%rax` (also known, for this problem, as `%eax`), so the result of `atoi` will be in `%rax`. After the call to `atoi` there is a `lea`, then an `add` and a `sub`. Where does the constant in the `sub` come from? .
16. Now you (kind of) understand the optimized `main`. What happened to the call to `loop_while`?
17. If you compile with `-O3` and disassemble the `main` program, you’ll discover that all traces of the loop have disappeared. We won’t try to analyze it — in particular, the instructions starting with `p` aren’t covered in this course — but it’s useful to know that the compiler has figured out the underlying math of `loop_while` and replaced it with a straight-line calculation. Wow!

Problem 2—Stepping and Looking at Data (19 Points)

Download the file `problem2.c` from the lab Web site. This file contains three `static` constants and three functions. Read the functions and figure out what they do. (If you’re new to C, you may need to consult your C book or some online references.) Here are some hints: `argv` is an array containing the strings that were passed to the program on the command line (or from gdb’s `run` command); `argc` is the number of arguments that were passed. By convention, `argv[0]` is the name of the program, so `argc` is always at least 1. The `malloc` line allocates a variable-sized array big enough to hold `argc` integers (which is slightly wasteful, since we only store `argc-1` integers there, but what the heck).

By now we hope you’ve learned that optimization is bad for debugging. So compile the program without optimization (but with `-g`!) and bring up the debugger on it.

1. gdb provides you lots of ways to look at memory. For example, type “`print puzzle1`” (something you should already be familiar with). What is printed?
2. Gee, that wasn’t very useful. Sometimes it’s worth trying different ways of exploring things. How about “`p/x puzzle1`”? What does that print? Is it more edifying?
3. You’ve just looked at `puzzle1` in decimal and hex. There’s also a way to treat it as a string, although the notation is a bit inconvenient. The “`x`” (examine) command lets you look at arbitrary memory in a variety of formats and notations. For example, “`x/bx`” examines bytes in hexadecimal. Let’s give that a try. Type “`x/4bx &puzzle1`” (the “`&`” symbol means “address of”; it’s necessary because the `x` command requires addresses rather than variable names). How does the output you see relate to the result of “`p/x puzzle1`”? (Incidentally, you can look at any arbitrary memory location with `x`, as in “`x/wx 0x8048500`”.)
4. OK, that was interesting and a bit weird (and we’ll be covering it in class soon). But we still don’t know what’s in `puzzle1`. We need help! And fortunately gdb has help built in. So type “`help x`”. Then experiment on `puzzle1` with various forms of the `x` command. For example, you might try “`x/16i &puzzle1`”. (`x/16i` is one of our favorite gdb commands—but since here we suspect that `puzzle1` is data, not instructions, the results might be interesting but probably not correct.) Keep experimenting until you find a sensible value for `puzzle1`. (Hint: Although `puzzle1` is declared as an `int`, it’s not. But on our machine an `int` is 4 bytes, 2 halfwords, or one—in gdb terms—word.) What is the human-friendly value of `puzzle1`? (Don’t accept an answer that is partially garbage!)
5. Having solved `puzzle1`, look at the value carefully. Is it correct? (You might wish to check it online.) If it’s wrong, why is it wrong?
6. Now we can move on to `puzzle2`. It pretends to be an *array* of `ints`, but you might suspect that it isn’t. Using your newfound skills, figure out what it is. (Hint: since there are two `ints`, the entire value occupies 8 bytes. What is the human-friendly value? It’s not “105”.)
7. Are you surprised?
8. Is it correct?
9. We have one puzzle left. By this point you may have already stumbled across its value. If not, figure it out; it’s often the case that in a debugger you need to make sense of apparently random data. What is stored in `puzzle3`?
10. We’ve done all this without actually running the program. But now it’s time to execute! Set a breakpoint in `fix_array`. Run the program with the arguments 1 1 2 3 5 8 13 21 44 65. When it stops, print `a_size` and verify that it is 10. Did you really need to use a `print` command to find the value of `a_size`? (**Hint:** look carefully at the output produced by gdb.)
11. *Existing breakpoint at fix_array.*
What is the value of `a`?
12. *Existing breakpoint at fix_array.*
Type “`display a`” to tell gdb that it should display `a` every time you stop. Step six times. You’ll note that one of the lines executed is a right curly brace; this is common when you’re in gdb and often indicates the end of a loop or the return from a function. After returning, what is the value of `a`?

13. *Existing breakpoint at fix_array.*

Step again (a seventh time). What is the value of `a` now? What is `i`?

14. *Existing breakpoint at fix_array.*

At this point you should (again) be at the call to `hmc_pomona_fix`. You already know what that function does, and stepping through it is a bit of a pain. The authors of debuggers are aware of that fact, and they always provide two ways to step line-by-line through a program. The one we've been using (`step`) is traditionally referred to as “step into”—if you are at the point of a function call, you move stepwise *into* the function being called. The alternative is “step over”—if you are at a normal line it operates just like `step`, but if you are at a function call it does the whole function just as if it were a single line. Let's try that now. In `gdb`, it's called `next` or just `n`. What line do we wind up at? (Incidentally, in `gdb` as in most debuggers, the line shown is the *next* line to be executed.)

15. *Existing breakpoint at fix_array.*

Use `n` to step past that line, verifying that it works just like `s` when you're not at a function call. What's `a` now?

16. *Existing breakpoint at fix_array.*

It's often useful to be able to follow pointers. `gdb` is unusually smart in this respect; you can type complicated expressions like `p *a.b->c[i].d->e`. Here, we have kind of lost track of `a`, and we just want to know what it's pointing at. Type “`p *a`”. What do you get?

17. *Existing breakpoint at fix_array.*

Often when debugging, you know that you don't care about what happens in the next three or twelve lines. You could type “`s`” or “`n`” that many times, but we're computer scientists, and CS types sneer at work that computers could do for them—especially mentally taxing tasks like counting to twelve. So on a guess, type “`next 12`”. What line are you at?

18. *Existing breakpoint at fix_array.*

What is the value of `a` now?

19. *Existing breakpoint at fix_array.*

What is the value of `*a`?

Finally, a small side comment: if you've set up a lot of `display` commands and want to get rid of some of them, investigate `info display` and `undisplay`.

Problem 3—Assembly-Level Debugging (19 Points)

So far, we've been taking advantage of the fact that `gdb` understands your program at the source level: it knows about strings, source lines, call chains, and even complicated C++ data structures. But sometimes it's necessary to get down and dirty with the assembly code.

To be sure we're all on the same page, let's quit `gdb` and bring it up on `problem2` again. Run the program with arguments of `1 42 2 47 3`.

1. What is the output?

2. Set a breakpoint in `main`. Run the program again. Where does it stop?
3. *Existing breakpoint at `main`.*
Type “list” to see what’s nearby, then type “b 35” and “c”. Where does it stop now?
4. *Existing breakpoints at `main` lines 33 and 35.*
So since that’s the start of the loop, typing “c” will take you to the next iteration, right?
5. *Existing breakpoints at `main` lines 33 and 35.*
Oops. Good thing we can start over by just typing “r”. Continue past that first breakpoint to the second one, which is what we care about. But why, if we’re in the `for` statement, didn’t it stop the second time? Type “info b” (or “info breakpoints” for the terminally verbose). Lots of good stuff there. The important thing is in the “address” column. Take note of the address given for breakpoint 2, and then type “disassem main”. You’ll note that there’s a helpful little arrow right at breakpoint 2’s address, since that’s the instruction we’re about to execute. Looking back at the corresponding source code, what part of the `for` statement does this assembly code correspond to?
6. *Existing breakpoints at `main` lines 33 and 35.*
The code at +44 jumps to `main+104`, which has three instructions that jump back to `main+46`. This is all part of the `for` loop pattern we covered in class. We’ve successfully broken (“broken?” “Set a breakpoint?”) at the initialization of the loop. But we’d like to have a breakpoint *inside* the `for` loop, so we could stop on every iteration. The jump to `main+46` tells us that we want to stop there. But that’s not a source line; it’s in the middle clause of the `for` statement. No worries, though, because `gdb` will let us set a breakpoint on *any* instruction even if it’s in the middle of a statement. Just type “b *(main+46)” or “b *0x40069f” (assuming that’s the address of `main+46`, as it was when we wrote these instructions). The asterisk tells `gdb` to interpret the rest of the command as an address in memory, as opposed to a line number in the source code. What does “info b” tell you about the line number you chose? (Fine, we could have just set a breakpoint at that line. But there are more complicated situations where there isn’t a simple line number, so it’s still useful to know about the asterisk.)
7. *Existing breakpoints at `main` lines 33 and 35, and instruction `main+46`.*
We can look at the current value of the array by typing “p array[0]@argc”. But the current value isn’t interesting. Let’s continue a few times and see what it looks like then. Typing “c” over and over is tedious (especially if you need to do it 10,000 times!) so let’s continue to breakpoint 3 and then try “c 4”. What are the full contents of `array`?
8. *Existing breakpoints at `main` lines 33 and 35, and instruction `main+46`.*
Perhaps we wish we had done “c 3” instead of “c 4”. We can rerun the program, but we really don’t need all the breakpoints; we’re only working with breakpoint 3. Type “info b” to find out what’s going on right now. Then use “d 1” or “delete 1” to completely get rid of breakpoint 1. But maybe breakpoint 2 will be useful in the future, so type “disable 2”. Use “info b” to verify that it’s no longer enabled (“Enb”). Continue past breakpoint 1, where we’re stopped. Where do we stop next? (Hopefully that wasn’t too much of a surprise!)
9. Sometimes, instead of stepping through a program line by line, we want to see what the individual instructions do. Of course, instructions manipulate registers. Quit `gdb` and restart it, setting a breakpoint in `fix_array`. Run the program with arguments of 1 42 2 47 3. Type “info registers” to see

all the processor registers in both hex and decimal. Which of the registers have *not* been covered in class?

10. *Existing breakpoint at fix_array.*

Well, that's because they're weird and not terribly important. (Except `eflags`, which holds the condition codes among other things. Note that instead of being given in decimal, it's given symbolically—things like `CF`, `ZF`, etc.) What flags are set right now?

11. *Existing breakpoint at fix_array.*

Often, looking at *all* the registers is excessive. Perhaps we only care about one. Type `"p $rax"`. What is the value? Is `"p/x $rax"` more meaningful?

12. *Existing breakpoint at fix_array.*

We mentioned a fondness for `"x/16i"`. Actually, what we really like is `"x/16i $rip"`. What do you see? Compare that to the result of `"disassem fix_array"`.

13. *Existing breakpoint at fix_array.*

Finally, we mentioned stepping by instructions. That's done with `"stepi"` ("step one instruction"). Type that now, and note that `gdb` gives a new instruction address but still says that you're in the `for` loop. Hit return to `stepi` again, and keep hitting return until the displayed line doesn't contain a hexadecimal instruction address. Where are you?

14. *Existing breakpoint at fix_array.*

It's useful to use `"x/16i $rip"` here to make sure we understand what's about to happen. You should see three `mov` instructions followed by a `call`. Use `stepi 3` to get past the `movs`. What instruction address will be executed next?

15. *Existing breakpoint at fix_array.*

As with source-level debugging, at the assembly level it's often useful to skip over function calls. At this point you have a choice of typing `"stepi"` or `"nexti"`. If you type `"stepi"`, what do you expect the next instruction to be (hexadecimal address)? What about `"nexti"`? (By now, your debugging `gdb` skills should be strong enough that you can try one, restart the program, and try the other, so there's little excuse for getting this one wrong!)

16. *Existing breakpoint at fix_array.*

Almost there! Stepping one instruction at a time can be tedious. You can always use `"stepi n"` to zip past a bunch, but when you're dealing with loops and conditionals it can be hard to decide whether it's going to be 1,042 or 47,093 instructions before you reach the next interesting point in your program. Sure, you could set a breakpoint at the next suspect line. But sometimes the definition of "interesting" is *inside* a line. Let's say, just for the sake of argument, that you are interested in how the `leavq` instruction works. You can set a breakpoint there by typing `"b *0x40066f"` (assuming that `0x40066f` is its address, as it was when we wrote these instructions). Do so, and then continue. What source line is listed?

17. *Existing breakpoints at fix_array and *0x40066f.*

The `leavq` instruction manipulates registers in some fashion. Start by looking at what `%rsp` points to. You can find out the address with `"p $rsp"` and then use the `x` command, or you could just try `"p/x $rsp"`. What are the values of `rsp` and `rbp`?

18. *Existing breakpoints at `fix_array` and `*0x40066f`.*

Use “`info reg`” to find out what all the registers. Then continue until you hit the `leave` instruction, and look at all the registers again. Have the values in the `rsp` or `rbp` registers changed, and what are their old and new values?